

# Midterm Review



Where have we been

ECE 373

# Prelims

- Questions on lab, homework or class?
  - Memory types
  - PCI connections
  - Debugging tricks
  - Delays and timers



# Types of OS's in the wild

- Single-user (Phones, PC's)
- Multi-user (Servers, mainframes, "cloud")
- Real-time (Stop lights, shuttle navigation)
- Embedded (Watch, routers, car engine, mp3)



# Device drivers

- The software that controls specific pieces of hardware
- API in the OS allows common interfaces for drivers
- Driver takes common commands from the OS and translates into hardware-specific stuff
- Many types of device drivers



# Types of drivers

- Three main classes of drivers



# Types of drivers

- Three main classes of drivers
  - Block drivers
  - Char drivers
  - Network drivers



# Types of drivers

- Three main classes of drivers
  - Block drivers
  - Char drivers
  - Network drivers
- Driver class can apply to many types of devices
  - USB device can be char, block, or network
  - Interface cards (PCIe) can be network or block
  - Graphics typically char devices

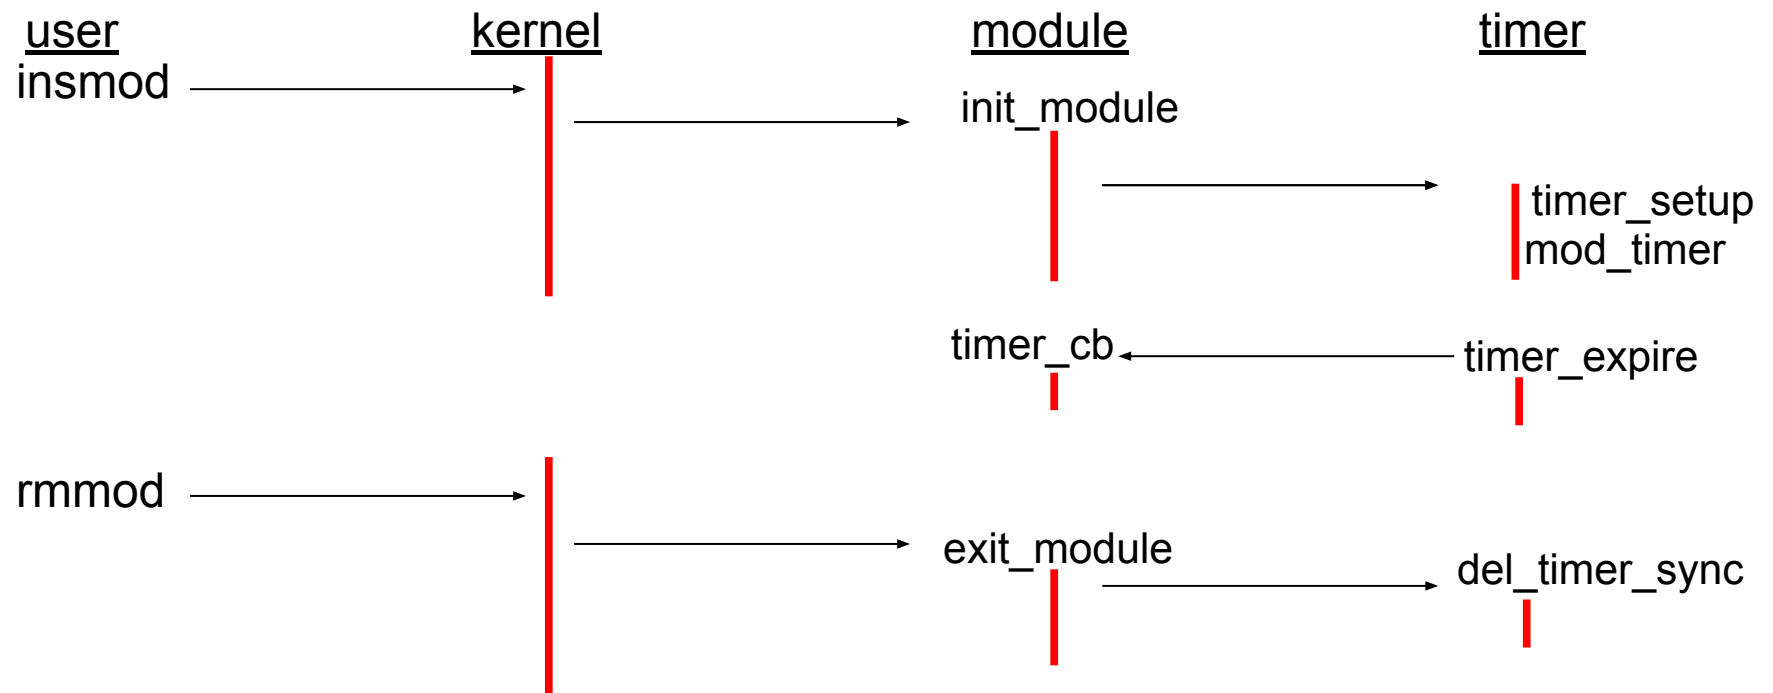


# Basic module concept



# Basic module concept

- Dynamic loading (insmod/modprobe, rmmod)
- Registering init and exit
- Kernel timeline of module diagram



# Bits and pieces

- Minimal callback hooks and Compile headers
  - Basic #includes
  - MODULE\_LICENSE(lic) – legal strings
  - \_\_init, \_\_exit
  - module\_init(func), module\_exit(func)
- Time values
  - Jiffies, HZ



# Module parameters



# Module parameters

- Gets info into the module
- `module_param(var, type, 0);`
- `/sys/module/<drivername>/parameters/<var>`
- **Kernel community typically frowns on them**
- `insmod <modulename> xx=11`
- `modinfo <modulename>`



# Talking to the Kernel



# Talking to the Kernel

- `ioctl()`
  - Older method of system call
  - Usually frowned up now-a-days (why?)
- Netlink
  - Newer, more flexible
  - Usually used with networking tools
- Pseudo filesystems
  - `/proc` – lots of kernel data
  - `/sys` – mostly module information
  - `/dev` – file-based connectors for read/write access
  - `debugfs` – useful for exposing debug hooks



# Character drivers



# Character drivers

- Operate on low-bandwidth devices
- Standard callbacks
  - Open, read, write, release
- Coordinate with OS
  - `alloc_chrdev_region()`
  - `cdev`
  - `struct file_operations`
  - Linkage into `/dev`





# I/O Ports



- 

- 

- 

- 

- 

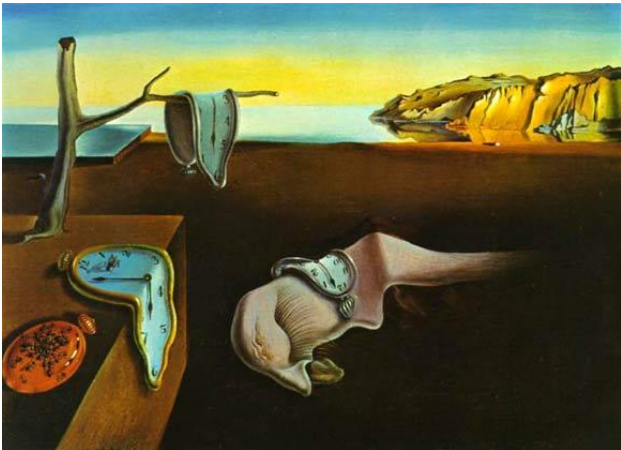
-

# I/O Ports




- Working with slow devices
- Different "address" space from system memory
- System design coordinates device address mapping
- PC architecture evolution
  - combined many little devices into fewer large devices
  - Adds complexity to use of devices

# Memory Map Types



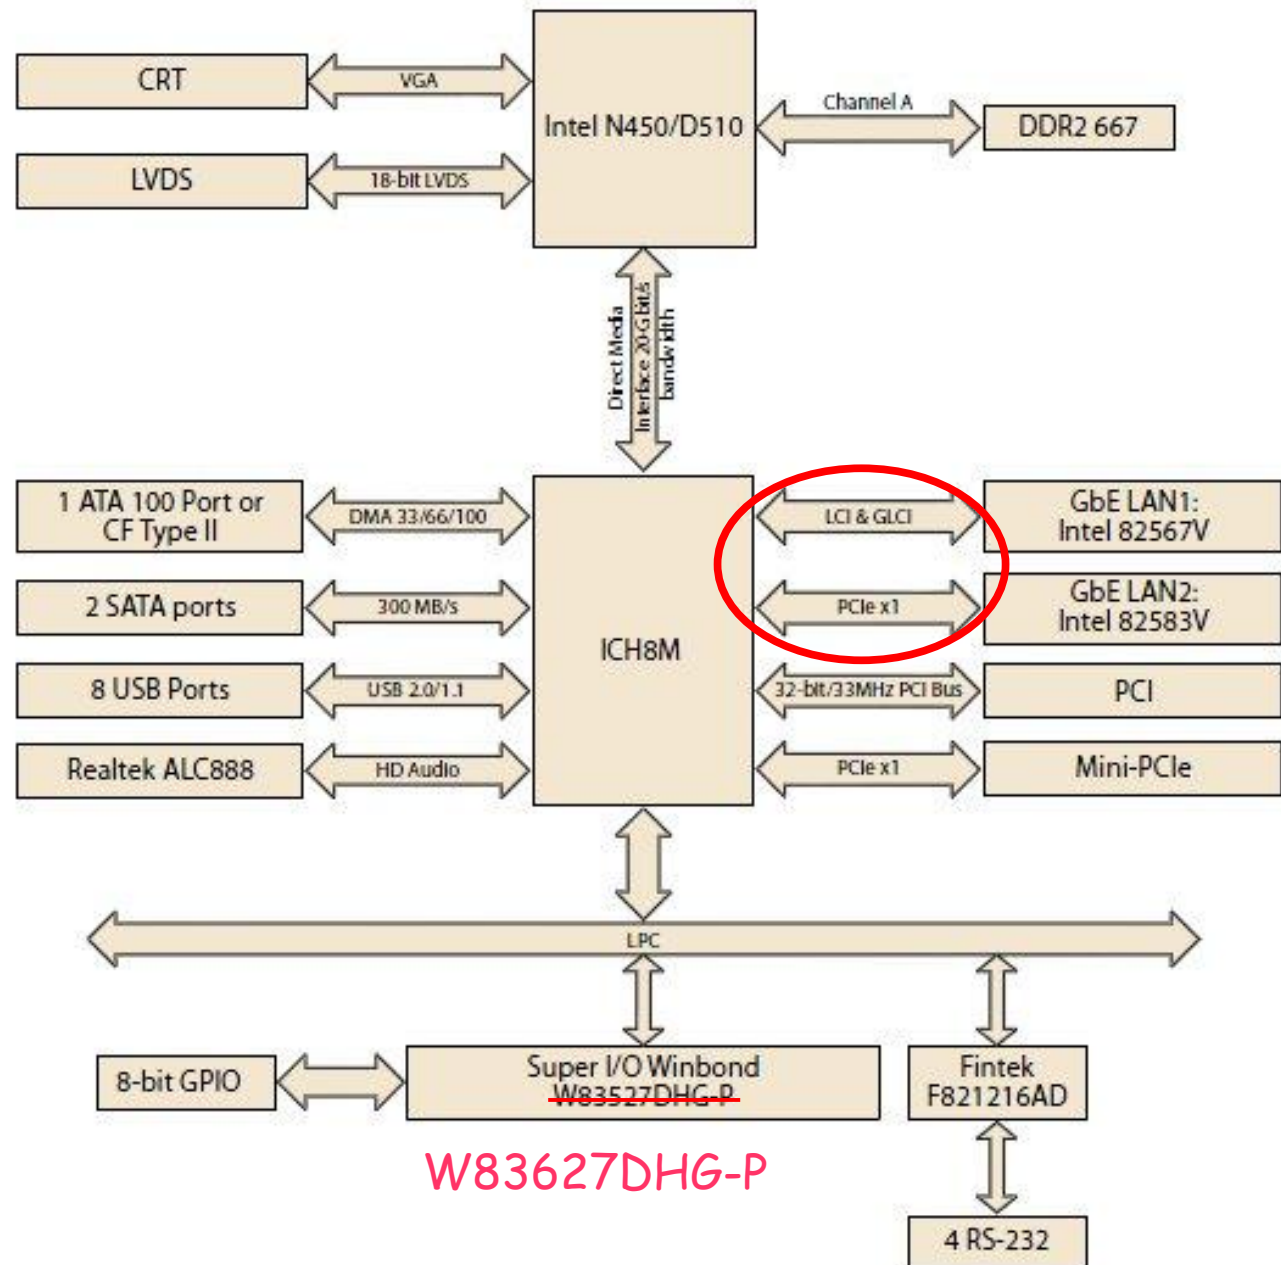
# Memory Types

- Physical addresses – actual memory address
- Virtual addresses – relocatable, flexible, swappable
- Kernel logical – maps directly to physical, used for DMA/sharing with devices
- Kernel virtual – relocatable kernel memory

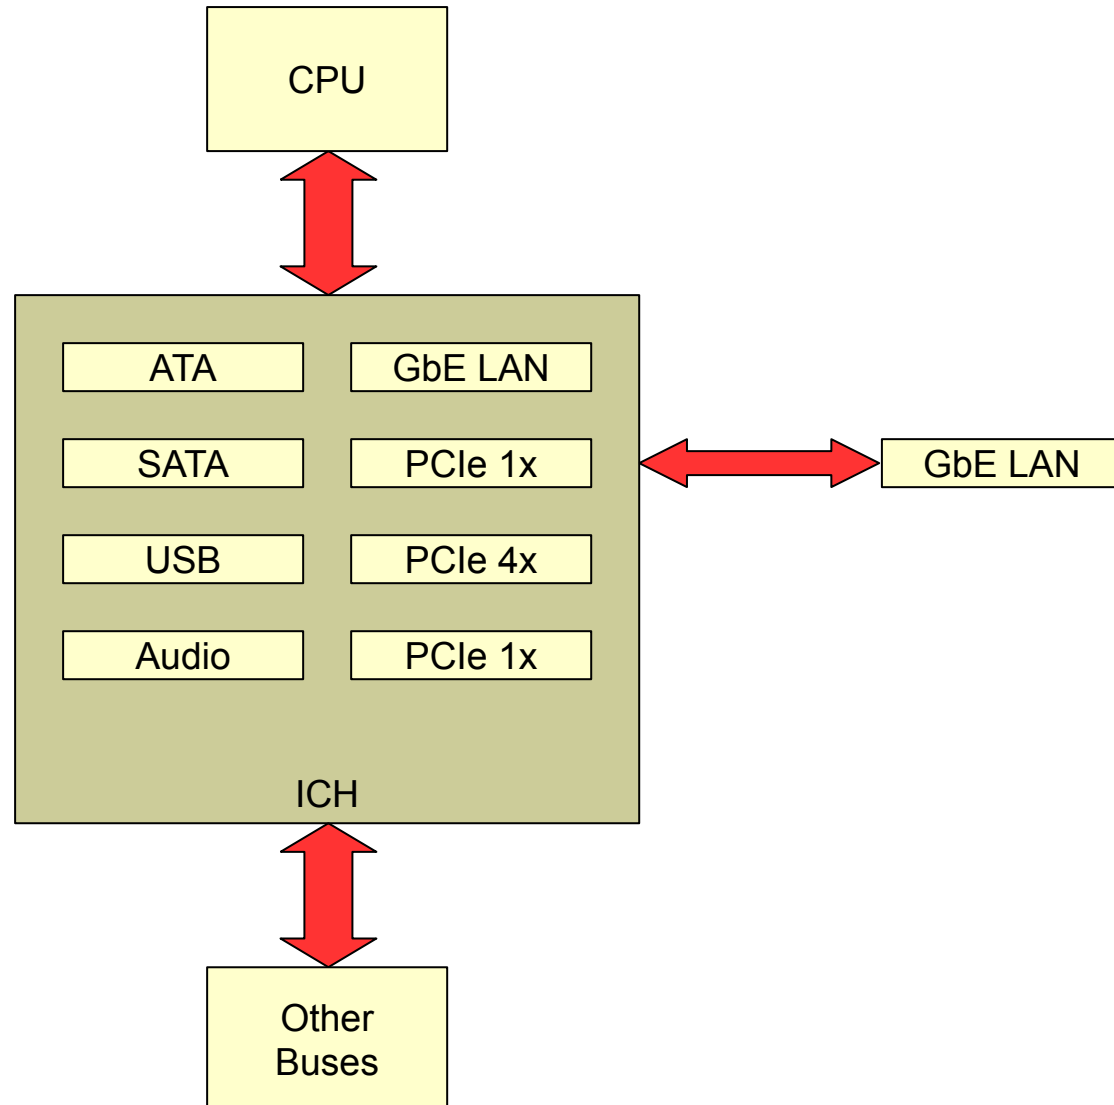


# Current PCs

- CPU
- Chipset
- Super I/O



# ICH Reality



# PCI



# PCI



- Communication/connection method for devices
- Devices use both port mapped and memory mapped I/O (focus on MMIO only for midterm)
- BAR – base address register
  - Starting address for device memory map
  - Driver uses `iowrite32()` and `ioread32()` to access device registers (or `writel()` / `readl()`)
- Callbacks
  - Probe, remove, sleep, resume



# Debugging

- Find a way to reproduce the bug
  - Gather info for tracking the problem
  - Prove it is fixed



# Debugging

- Find a way to reproduce the bug
  - Gather info for tracking the problem
  - Prove it is fixed
- Printk
  - KERN\_INFO, etc. Where do these go?
  - WARN(), BUG()



# Debugging

- Find a way to reproduce the bug
  - Gather info for tracking the problem
  - Prove it is fixed
- Printk
  - KERN\_INFO, etc. Where do these go?
  - WARN(), BUG()
- Stack trace
  - Gives idea of where things broke
  - Objdump helps decode stack info



# Debugging

- Find a way to reproduce the bug
  - Gather info for tracking the problem
  - Prove it is fixed
- Printk
  - KERN\_INFO, etc. Where do these go?
  - WARN(), BUG()
- Stack trace
  - Gives idea of where things broke
  - Objdump helps decode stack info
- /proc – various kernel data for monitoring
- Kgdb – useful to break into kernel
- Ethtool – monitor network devices
- Bus trace tools



# Delays and time keeping

- How jiffies relate to time
- Delays vs. Sleeping
- Pre-emption and its impacts
- `get_cycles()` and why use it



# Time for Timer

- Basics

- `timer_setup(t_var, t_callback, t_flags)`
- `mod_timer(t_var, interval)`
- `del_timer_sync(t_var)`



# Time for Timer

- Basics

- `timer_setup(t_var, t_callback, t_flags)`
- `mod_timer(t_var, interval)`
- `del_timer_sync(t_var)`

- Callback function

- Called when timer expires
- Given pointer to `t_var` as an argument
- `timer_cb(struct timer_list *t)`
- Use `from_timer()` to get back to private driver struct



# Getting information

- Device spec sheets
- The web
- Books
- Kernel and driver source code

